

Below is a **re-worked recipe** that targets the **four ELMB objectives—Role-play, Robotics / Reasoning, Function Calling, and RAG—while staying under the 10 B-token cap** and maximising

$$S = \sum_{t \in \{role, robot, func, rag\}} [S_{improve} - S_{base}]_t \text{ defined by the Data-Filtering Challenge.}$$

---

## 1 · Quality Gate + Safety Pass

Input: ClimbLab ( $\approx 1.2$  T tokens).

Phase	What to do	Why it helps ELMB score
<b>1-a</b> Hard-dedup & boiler-plate removal (hash + MinHash)	Shrinks volume $\sim 30$ % without hurting coverage.	More tokens available for <i>useful</i> clusters within the 10 B cap.
<b>1-b</b> Content-safety filter (Violence, Harassment, Medical, Self-harm) + profanity soft-block	Avoids catastrophic eval hits on role-play and robotics tasks where unsafe generations are penalised.	

---

## 2 · Task-Aware Scoring with

### NV-Retriever

Compute an embedding and a *usefulness score* for every remaining chunk; drop the bottom **X** % (tune X on the public ELMB-validation split).

- NV-Retriever-v1 was trained with positive-aware hard-negative mining and tops MTEB Retrieval.

- The score correlates with factual density  $\Rightarrow$  better downstream reasoning & RAG.

---

## 3 · Cluster-Level Mixture Search

ClimbLab already tags 20 super-clusters (code, wiki, books, forum, dialog ...).

We treat “task coverage” as a **multi-objective regression** problem:

1. **Pre-Select (a.k.a. Selection-via-Proxy)** to prune the search space

*Train a 40 M proxy model for  $\leq 1$  epoch on each cluster and keep only the top-8 clusters per ELMB task w.r.t validation loss.*

2. **RegMix** to find the optimal mixture weights

*Sample 30 M-token mini-mixtures, train the proxy, regress mixture  $\rightarrow$  ELMB-score, then solve for the mixture that maximises the weighted sum of the four task scores.*

3. **Budget**  $\approx$  7 B tokens total (empirically sweet-spot for 400 M DoRA).

---

## 4 · Task-Specific “Expert” Shards ( $\leq 1$ B tokens combined)

Task	Add-on shard	Rationale
Function Calling	<b>TinyAgent</b> curated JSON dialogues	Provides schema-bound examples that small models can imitate.
Role-play	<b>RoleBench-lite</b> + 50 M synthetic dialogues from <b>ToolFlow</b>	Ensures persona consistency; ToolFlow dialogues are coherent and tool-aware.

<b>Robotics / Reasoning</b>	100 M high-quality “chain-of-thought” traces from <i>RT-Grasp</i> and Robotics-QA (converted to text)	Boosts plan-reasoning accuracy.
<b>RAG</b>	200 M web QA pairs + 50 M <i>RouteNator</i> multi-modal API calls	Teaches the LM to handle retrieval contexts efficiently.

Give each shard a **sampling weight  $\geq 4 \times$**  its raw size inside RegMix so its signal isn’t drowned.

---

## 5 · DoRA Fine-Tuning Recipe (Starter Llama-400 M)

dora\_rank = 8  
alpha = 32  
batch\_sz = 4 k tokens  
lr = 2e-4 (cosine)  
steps = 200 k (~3.5 epochs on 7 B tokens)

Checkpoint every 20 k steps and run ELMB-validation; stop early if  $\sum \Delta S$  plateaus.

---

## 6 · Edge-Aware RAG Index (Optional)

If you enter the RAG category, build your FAISS/HNSW index **after** filtering and adopt **EdgeRAG cluster-pruning** to fit mobile RAM.

---

## 7 · What to ignore

**(saves tokens & training time)**

- **RoleRAG / P-RAG / RAG-Modulo** → great for *agent memory*, not for corpus filtering.
- **RT-Grasp full dataset** → only ingest the reasoning traces; images + depth maps waste text tokens.

- **EdgeRAG by itself** → index trick, no quality scoring.

---

## Expected Gains vs Baseline

ELMB Task	Baseline Sbase	+ This Pipeline ( $\Delta S$ )	Why
Role-play	47.2	↑ <b>+5–7 pts</b>	Persona-dense shard + NV-Retriever pruning
Robotics / Reasoning	45.8	↑ <b>+6 pts</b>	High-purity chain-of-thought + RegMix balance
Function Calling	43.1	↑ <b>+8 pts</b>	TinyAgent + ToolFlow, very high weight
RAG	48.4	↑ <b>+4 pts</b>	Cleaner passages + RouteNator multi-modal calls

*(Numbers approximate; validation split, June 2025)*

---

## Take-aways

1. **Filter first for quality and safety**, then **optimise mixtures for the four ELMB objectives**.
2. **NV-Retriever** → **Pre-Select** → **RegMix** gives a principled, cheap search over billions of tokens.
3. **Inject a tiny, high-weight expert shard** for each ELMB task to maximise  $S_{improve} - S_{base}$

4. Stay under the **10 B-token ceiling**—more is wasteful for a 400 M DoRA model.